Intelligent Extraction. Deterministic Results.

Solving the Insurance Comparison Challenge by keeping LLMs away from the math.

Probabilistic Intelligence

Gemini 2.5 Flash

Used EXCLUSIVELY to read, parse, and structure messy RFP inputs across segments. The LLM never makes decisions.

Deterministic Logic

Python Win-Count Logic

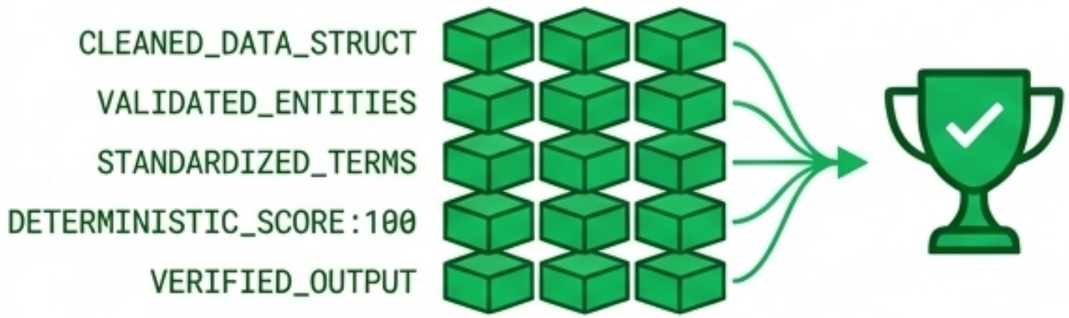

Strict Python win-count logic. The code dictates the outcome. The LLM is entirely locked out of the decision-making process.

FastAPI | Google Cloud Run | PostgreSQL Caching

NotebookLM

The MrRooT.Ai Processing Pipeline

A production-ready architecture secured by a strict "Critical Boundary".

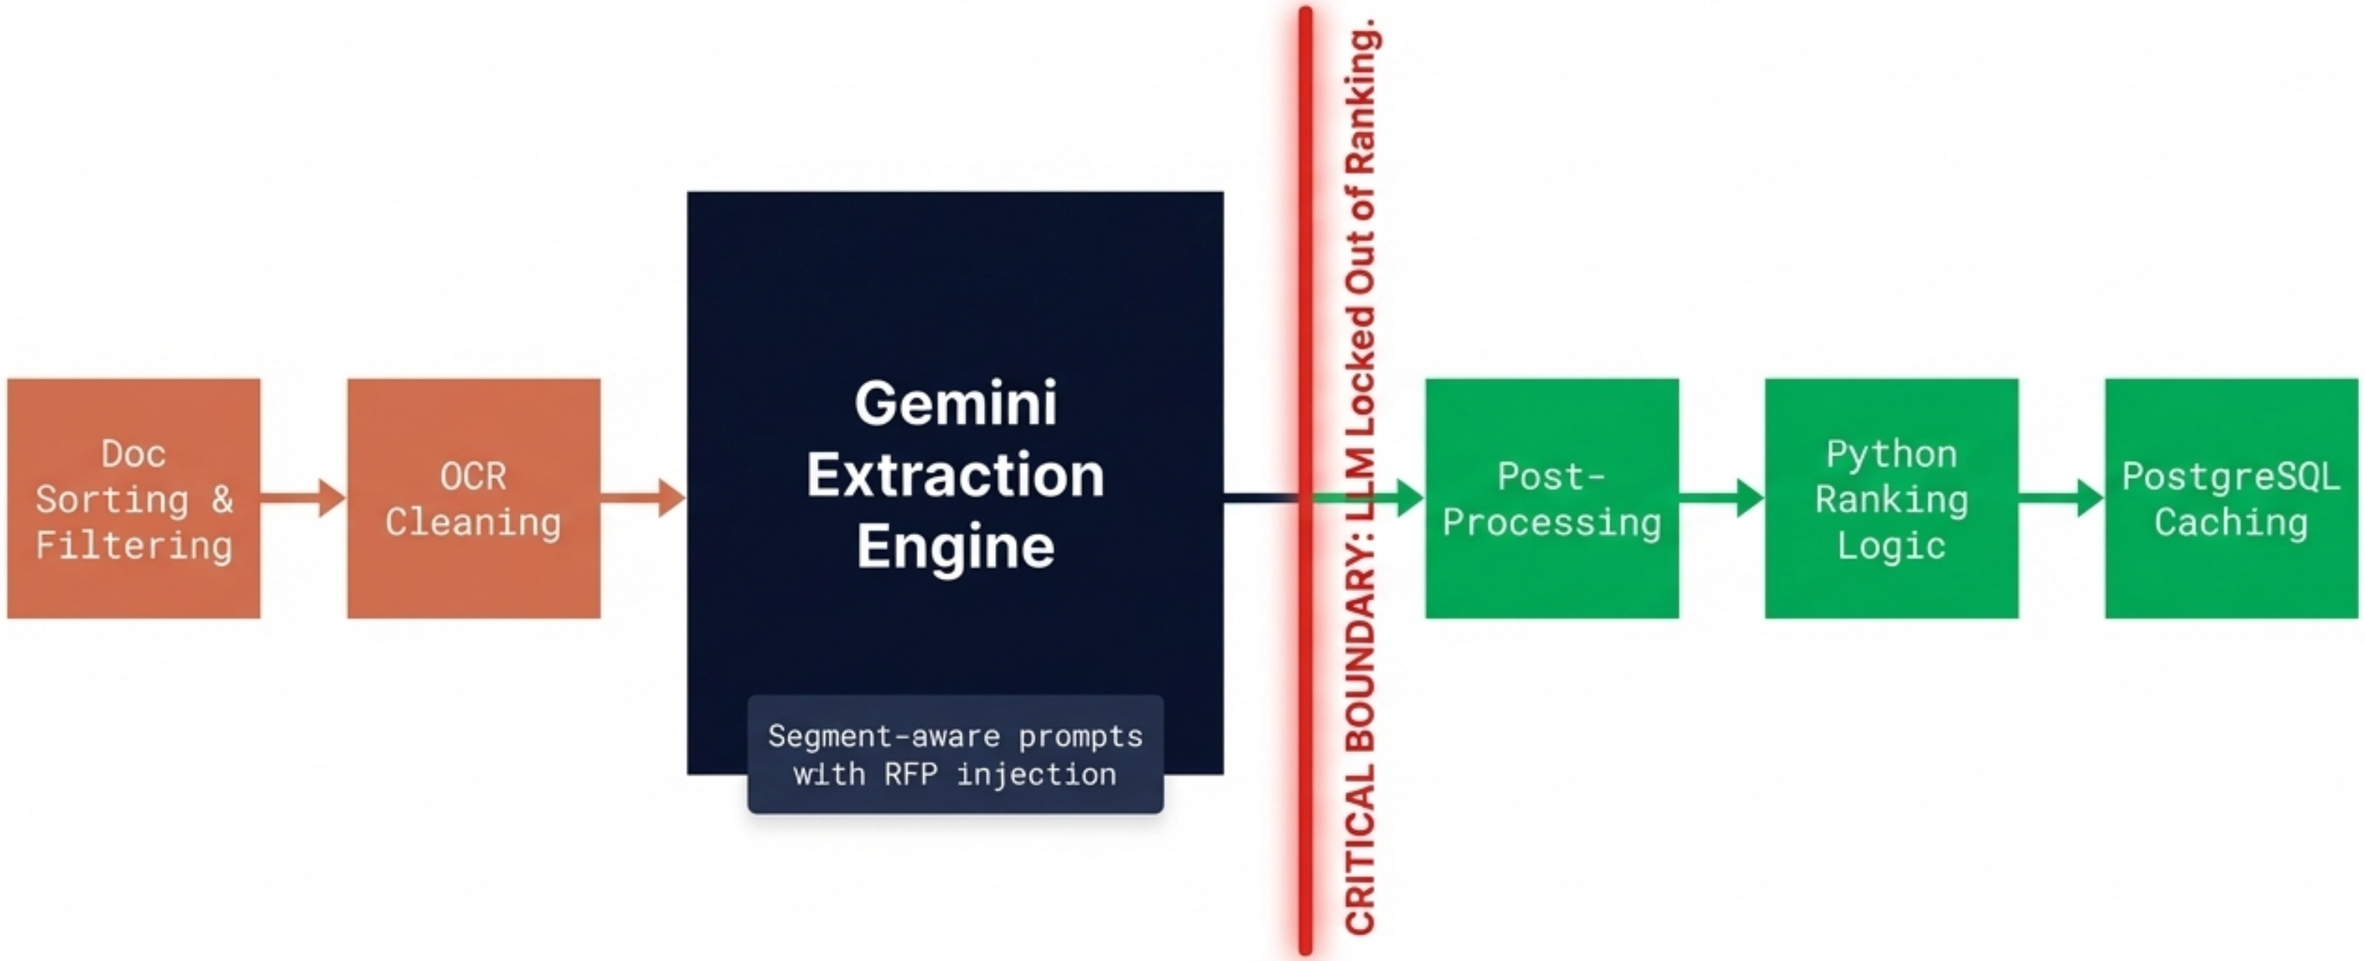

NotebookLM

Performance Diagnostic: Isolating the Variance

Flawless logic, bottlenecked strictly by baseline OCR.

Overall: 0.589

Segment	Extraction	Ranking	Best Offer	Total
auta	Extraction: 0.754	Ranking: 1.000	Best 1.000	0.853
lodě	Extraction: 0.160	Ranking: 1.000	Best 1.000	0.496
odpovědnost	Extraction: 0.136	Ranking: 0.750	Best 1.000	0.419

1.000

Best Offer Selection Accuracy across ALL three segments.

Anatomy of an Extraction Failure.

Limit I	Limit II
50,000 CZK	100,000 CZK

Standard Base OCR

Limit I
Limit II
50,000 CZK
100,000 CZK

Main Bottleneck: Baseline OCR destroys the spatial reality of complex liability tables before the data ever reaches the Gemini prompt.

The logic engine is highly stable. The path to 0.85+ overall requires zero logic changes.

Next Step: Table-Aware Parsing / Re-OCR

Integrating spatial OCR preserves column alignments for Odpovědnost, instantly unlocking the remaining fields.

NotebookLM